

A Report Of “EMERGE: Enhancing Multimodal Electronic Health Records Predictive Modeling with Retrieval-Augmented Generation”

Guoqian Zeng¹, Paul Grigas¹, Qiran Dong¹
j_zeng@berkeley.edu, pgrigas@berkeley.edu, qirandong@berkeley.edu

1: University of California, Berkeley, Berkeley, California.

I. Summary of Contributions

The public healthcare infrastructures and system has a high demand on intelligent renewing of the delivering system, like staff planning or readmission risk management, to improve treatment success rate and operational efficiency. Operational work in the health system produces vast, continuous, and heterogeneous records that could be used in multimodal predictive learning. The emerging Electronic Health Records (EHRs) consolidate digitized health information from multiple sources, organized around the patient, often containing a patient’s medical history, diagnoses, immunization dates, and clinical notes. Such data type was used in previous studies with neural network approaches singular-model predictions. Existing multimodal models process to learn a mapping from heterogeneous inputs to output labels on other data types. In concern to the lack of medical expertise to weigh health information, some recent studies have begun to integrate knowledge graphs (KG) in their analysis.

However, the paper notices a gap between the EHRs’ abundant source of information—especially previously ignored time-series data and text-free clinical notes—and the effectiveness of knowledge graph insertion for prediction modeling.

Therefore, a predictive modeling framework with integrated RAG, EMERGE, was proposed by the paper as a modern approach to improve effectiveness of healthcare predictive tasks with EHRs data. For this report, we divided it into two-stages: (1) an extra summary generation driven by a RAG subframework; followed by (2) a predictive multimodal modeling.

EMERGE is the state-of-the-art framework that tackles these challenges to elevate predictive modeling in EHRs data, with its combined structure of the Retrieval Augmented Generation and the multimodal fusion. This structure helps it utilize a broad range of professional information, while connecting them in a biomedical perspective, to elaborate prompts for health specific predictive tasks.

To ensure the goal-oriented feature, the paper further concludes two major challenges to focus while developing solutions: 1. How to perform entity extraction and matching from EHR analysis to incorporate with KG? 2. What helps integrating long-text retrieval for task-specific prompting?

Initially, they want to get the representation from the original data as a part of input to the LLM before using LLM as an enhancement, this would mine the meaningful information from the original EHRs data to further improve the effectiveness of their innovative models. They use Gated Recurrent Unit network as the encoder to capture the time-dependencies in sequence data and encode this temporally linked information. They then utilize a medical domain language model to obtain text embeddings.

Let's now focus on the Retrieval Augmented Generation part, this part extracts medical entities from time-series data and texts information from the EHRs to generate a concise health summary. Purposes of generating the summary is to use as supplemental information like an extension with the mixed of the original information for the BERT model in the later part to perform specific predictive tasks. Why do we need these extension data? Because it would provide the accurate information, specific sourced information that is cited directly from the patient records, and overcome terminology confusion to provide domain-specific results for the task, accurating predicting. (Why do they use text-to-category? Are there any other expected outcomes they described?)

In terms of time-series data extraction, they use a z-score method, which is also widely being used in other medical fields, to find the outliers of the data. The outliers represent the abnormal high or low in the tabulated data of measurement, such as the temperature is too low. These outliers are significant to identify the patient disease or parts to be noticed.

Text extraction, they use an iterative method to first use LLM to conclude the important disease words in the clinical notes, and then do a refinement to improve quality. In the first stage, they designed a prompt with clear instructions and an example to tell the LLM to extract disease terminologies that the patient might suffer from. They then use the prompt to generate information, extended with the original inputs. In refinement, they deal with hallucinations by three steps: filter out the terminologies that are not in the original input, filter out the words that are not in their defined disease type, and filter out duplicate words to avoid semantic redundancy. They replicate this process for multiple times until achieving convergence.

Now, after getting the entities as extraction of time-series data and clinical notes, the paper goes on to use a semantic-based vector retrieval approach to tackle the problems in matching of medical entities and knowledge graphs. Initially, they used the same sentence embedding model to encode nodes from Knowledge Graph and the entities from EHRs, ensuring that these embeddings are aligned with the same vector space. Then, for the matching process, they calculate the similarities between the entity and each node in the Knowledge graph using the cosine similarities. They published a threshold along with their methods to select the matched nodes while maintaining qualities. This means that they will search for the most nodes with the highest similarity with the KG node and filter out the maximums below a certain level. After the

framework selects a set of matched nodes, it gets definitions and descriptions of entities from the KG. Later it obtains the relationships between diseases which act as edges of KG.

The framework walks into the procedure to generate a prompt for the concise patient health summary with the information entities drawing from EHR time-series data and clinical notes with their extended information gained. The framework uses RAG to dense information because a combination of complete information of extracted entities from original data set and supplementary information might be overlong for the input of LLM. A prompt template consisting of (1) role & instructure (2) extracted entities in time series (3) extracted entities in notes (4) retrieved knowledge graph nodes (5) retrieved knowledge graph triplets is provided.

Now, with extracted entities from RAG and separate extracted entities from clinical text and time-series data, the Multimodal Fusion Network is structured to integrate these information in an unified version and fuse for prediction. The embeddings from clinical notes and RAG summaries, both texts, are concatenated and passed through a text fusion module to generate a combined text representation. Then, a cross-attention mechanism is employed in both directions: the time-series representation attends to the fused text representation, and vice versa. This bidirectional attention allows each modality to incorporate contextual information from the other. The outputs from both attention pathways are then concatenated and passed through a Multi-Layer Perceptron (MLP) to get the final fused representation for the prediction tasks. The BCE loss is selected for downstream prediction tasks such as in-hospital mortality and 30-day readmission.

This structure is strongly correlated with our course because it elevates the LLM for an engineering design to assist the healthcare operations. This is related to the time-series modeling in 242A and the recurrent neural network modeling in 242B. We are able to use the understanding of the neural network architectures to improve this study.

II. Evaluation of the Contributions

The EMERGE framework demonstrates its state-of-the-art performance as an integration of RAG with multimodal EHR data to enhance predictive modeling in the healthcare delivery system. This evaluation critically assesses the strengths and limitations of EMERGE, providing insights into its contributions and areas for improvement.

1. Strengths

It innovates RAG integration on LLMs with both time-series data and clinical notes to extract and summarize relevant medical entities. By aligning extracted entities with a professional medical knowledge graph (PrimeKG), EMERGE ensures consistency and semantic richness in the data representation. This approach enables model learning to absorb multimodal data with contextual information and capture professional relationships between entities, which is a potential cause for the outperformed model accuracy in the experiments.

When bridge the extractions from multimodal resources, the framework employs an adaptive cross-attention mechanism to fuse information on an uniform dimension effectively. This design allows for a more comprehensive understanding of patient pattern, as it considers the interplay between various data types.

In experiments, it demonstrates superior prediction and AUROC performance on benchmark datasets (MIMIC-III and MIMIC-IV) for in-hospital mortality and 30-day readmission prediction. Then, it also demonstrates the workflow design process with performance comparison of critical points like GRU and adaptive features with their counterparts. The inclusion of extensive ablation studies further validates the effectiveness of each component within the framework.

However, EMERGE uses complex integration of multiple components, including LLMs, knowledge graphs, and cross-attention mechanisms, which increases the computational demands of the framework. This complexity may hinder real-time application in clinical settings where resources are limited. Therefore, streamlining the framework through model pruning or more efficient algorithms could possibly reduce computational overhead.

Notably, its reliance on external knowledge graphs like PrimeKG introduces potential challenges related to data availability and maintenance. Any inconsistencies or updates in the knowledge base could affect the model's performance.

In the experiments, its validation is primarily based on MIMIC-III and MIMIC-IV datasets for in-hospital mortality and 30-day readmission prediction. Its performance and adaptability to other healthcare systems or diverse patient populations remain to be explored.

We must also recognize that the difficulties to apply dynamic knowledge graphs of this model would cause loss of model accuracy in clinical study, but developing advanced techniques to filter and prioritize retrieval could mitigate the risk of information loss and improve model interpretability.

III. Further Extensions

Building on EMERGE's RAG-driven fusion, we designed a natural extension of the framework which replaces its GRU and generic encoder components with a slightly more lightweight, modality-specific architecture that explicitly models time-series and text in parallel before fusing them. Specifically, this alternation employs a parallel processing—single-layer Feedforward Neural Network (FFN) to process normalized time-series features (e.g., vitals, labs) and a single-layer Recurrent Neural Network (RNN) to encode textual clinical notes—for different modalities. Research shows that even simple FFNs can match or outperform complex architectures on time-series tasks when appropriately configured, avoiding unnecessary complexity. Likewise, RNNs have demonstrated competitive performance on clinical event prediction from EHR sequences, indicating that a single-layer GRU suffices for extracting temporal patterns from text data.

To ensure each branch learns robust representations, our design suggests one can pretrain the FFN and RNN separately: first train the FFN on time-series-only readmission labels, then independently train the RNN on text-based predictive tasks such as readmission risk from discharge summaries. After such pretrain, the last hidden representations (h_t for the one of time-series-only training, h_c for the one of text-based training) are concatenated and passed through a linear fusion layer with learned weights to output final risk scores. This two-stage approach leverages modality-specific strengths and allows for targeted hyperparameter tuning per branch, following best practices in multimodal pipelines.

For evaluation of this new study, a focused ablation would quantify each modality's contribution by comparing:

- FFN only (time-series branch alone),
- RNN only (text branch alone),
- Combined FFN+RNN (our fusion model).

Such systematic ablation is critical to understand the utility of multi-modal integration. Prior work on medical time-series tasks underscores that ablation experiments reveal how each modality and fusion method affects overall performance, guiding efficient model design. Additionally, investigating fusion weights and branch-specific performance under data-scarcity scenarios can highlight robustness benefits, as demonstrated in irregular modality ablation studies.

Together, these extensions—single-layer FFN and RNN, separate pretraining, and a targeted ablation regime—offer a streamlined yet insightful advancement of EMERGE. Methodologically, this approach balances methodological rigor with computational feasibility, illuminating the distinct value each modality brings to readmission risk prediction. To achieve a more aligned comparison with EMERGE, we suggest further research could use RAG-driven integrations.

IV. References

- Chandak P, Huang K, Zitnik M. Building a knowledge graph to enable precision medicine. *Sci Data*. 2023 Feb 2;10(1):67. doi: 10.1038/s41597-023-01960-3. PMID: 36732524; PMCID: PMC9893183.
- Jha, A. K., DesRoches, C. M., Campbell, E. G., Donelan, K., Rao, S. R., Ferris, T. G., Shields, A., Rosenbaum, S., & Blumenthal, D. (2009). Use of electronic health records in U.S. hospitals. *New England Journal of Medicine*, 360(16), 1628–1638.
- Lewis, P., Oguz, B., Rinott, R., Riedel, S., & Stenetorp, P. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *Advances in Neural Information Processing Systems* 33 (NeurIPS 2020).
- Alvin Rajkomar, Eyal Oren, Kai Chen, Andrew M Dai, Nissan Hajaj, Michaela Hardt, Peter J Liu, Xiaobing Liu, Jake Marcus, Mimi Sun, et al. 2018. Scalable and accurate deep learning with electronic health records. *NPJ digital medicine* 1, 1 (2018), 18.
- Yinghao Zhu, Changyu Ren, Zixiang Wang, Xiaochen Zheng, Shiyun Xie, Junlan Feng, Xi Zhu, Zhoujun Li, Liantao Ma, and Chengwei Pan. 2024. EMERGE: Enhancing Multimodal Electronic Health Records Predictive Modeling with Retrieval-Augmented Generation. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management (CIKM '24)*, October 21–25, 2024, Boise, ID, USA. ACM, New York, NY, USA, 11 pages. DOI: 10.1145/3627673.3679582